

Pure Exploration

Tim Radtke

5/3/2017

In some previous chapter, we must have introduced the Canonical Exponential Family (see Garivier and Kaufmann, page 3) and the Kullback Leibler distance; especially also the KL for exponential family, for Bernoulli and Gaussian in particular, and so that $kl()$ is defined.

Before this chapter, a general introduction to A/B tests, with general sampling bounds for two variants case and uniform sampling. Get people hooked by general application in website optimization?

0.1 Multi Armed Bandit and Pure Exploration

We now move beyond the previously described problem of A/B testing by opening restrictions which we held so far. There are several ways in which to do that, but the main way will be two no longer require uniform sampling strategies. Further, we most often are interested in finding the best among K arms. Other parts of the literature try to find the best m among K arms.

(What should come first? Paper by Kaufmann et al. about A/B tests but unifying FC and FB, or first make a differentiation between FB and FC?)

For a start, consider again the idea of sampling K variants of for example websites which return a real-valued (as simple as Bernoulli) feedback coming from distributions ν_k , $k \in \{1, \dots, K\}$. We still want to find the distribution k for which the mean is equal to the highest mean among the distributions $\mu = \mu^*$. To do so, we are no longer restricted to sampling the distributions in a round-robin fashion as previously, but we are allowed to set up new sampling strategies with the general purpose of finding the best arm as quick as possible with as much confidence as possible.

To be a little more precise, we consider K probability distributions ν_1, \dots, ν_k with respective means μ_1, \dots, μ_k . At each round $t = 1, 2, \dots$, we can decide from which distribution to sample, that is, we pick an arm $A_t \in \mathcal{A} = \{1, \dots, k\}$, and we observe the feedback X_t from distribution μ_{A_t} . What distinguishes this from the A/B testing setting is the fact that by picking an arm we actively choose at each time the distribution from which we would like to receive the next sample. We are no longer restricted to the uniform sampling strategy. In fact, our goal is define the a sampling strategy $(A_t)_t$ with characteristics that are to be analyzed and ideally optimal in a sense to be specified.

In any case, the literature considered in this chapter is concerned with strategies whose objective it is to identify the best arm defined by $\mu^* = \max_{1 \leq a \leq K} \mu_a$. In contrast to multi-armed bandit problems concerned with optimizing regret (Auer et al, Agrawal et al), such

algorithm could in an extreme case sample the best arm only once but the inferior arms multiple times, if this helps identifying it as the clearly best arm for some reason. In general, it is only important to make a single correct decision *after* the final sample has been drawn.

As described in Garivier and Kaufmann (2016), a general strategy to find μ^* can be defined by three attributes:

(Add sigma field definition???)

- a *sampling rule* $(A_t)_t$
- a *stopping rule* τ
- a *decision rule* \hat{a}_t ($= \hat{a}_t(\delta)$)

This allows to describe the goal then as finding $\hat{a}_t \in \arg \max_a \mu_a$ with the largest confidence δ possible while constrained to minimizing the rounds τ needed to find the estimate and confidence.

Recent literature has dealt with this problem by specifying one of $\{\tau, \delta\}$ upfront and then optimizing the other. These two settings are called *fixed budget* (FB) and *fixed confidence* (FC).

Fixed confidence can seem like a more natural approach to the problems at first. Especially when considering the notion of statistical significance and the famous $\alpha = 0.05$ confidence level that comes with the attitude of “We do not accept any difference as significant unless a difference as large or larger as the one observed has a probability of less than 5% assuming there is no effect”. In the problem of identifying a superior arm, assuming there exists a superior arm, this notion might lead to very large sample sizes depending on the effect size. In particular, in the application of web optimization, where websites are optimized by conversion rate, which might be at a level of 0.06, for example, it would take 41575 samples to detect a 10% (or larger) difference in conversion rate at a 5% significance (and 95% power) [source](#).

But in the situation described above, we’re not really fixing the confidence level. We are fixing the number of samples to make sure that a given effect size could be detected as significant difference at some confidence level (here, 5%). But even if the actual difference between the two versions is less than the minimum detectable effect size, the test could give us some detected difference at some confidence. Looking at it this way, the fixed budget setting may actually be more alike to standard A/B testing procedures, in which a sample size is fixed *before* the test is run and determined by the minimum detectable effect, and confidence level and power of the test.

Consequently, the fixed confidence setting is concerned with fixing a maximal level of risk δ that one is willing to accept. While minimizing the expected sample size called *sample complexity*, a corresponding strategy needs to fulfill $\mathbb{P}(\hat{a}_\tau \notin \arg \max \mu_a) \leq \delta$ (and then is called δ -PAC). Compare Garivier and Kaufmann (2016).

On the other hand, the fixed budget setting describes strategies which fix the number of samples τ upfront and then try to minimize the probability of error $\mathbb{P}(\hat{a}_\tau \notin \arg \max \mu_a)$ of the decision that’s made after the last sample has been drawn.

Alternative formulations which are more rare in the literature are so called *anytime* strategies

aiming to optimize the probability of error and the sample complexity in a way such that the procedure can be stopped at anytime (after every round) and return a decision with a certain guaranteed level of risk. See [Jun and Nowak \(2016\)](#) for one discussion of the topic. The advantage over the fixed budget case is that users do not have to specify a budget in advance. This is especially helpful, if it is not foreseeable how many samples will be available. Since some fixed budget strategies explore arms in stages, they might not be able to guarantee a confidence when stopped before their budgeted rounds.

Other literature tries to analyze the two settings of fixed budget and fixed confidence separately, but to make them comparable by using complexity measures for problems that are on the same scale (see Kaufmann et al. 2016, Kaufmann et al. 2017). We will come back to this topic in chapter ...

0.1.1 Fixed Confidence

The setting of fixed confidence can be framed as a sequential testing problem. After each round, we might ask ourselves whether we have sufficient evidence to make a δ -PAC decision. If we cannot make such decision, we continue to draw the next sample. And because we are not restricted to uniform sampling strategies, we pull the next sample so as to gain as much information with the purpose of minimizing the expected overall number of samples until we can make the δ -PAC decision.

Descriptions like this lead back to Chernoff (1959) describing settings in which experiments (arms) are performed sequentially. After each round, it has to be decided whether a next experiment is performed or whether the final decision is made based on a hypothesis test where the decision selects between one of two actions. In contrast to the modern literature, Chernoff takes into account the cost that samples have and acknowledges that his derivations hold when the cost of sampling goes to zero. In particular, he compares two Bernoulli distributions based on their means μ_1, μ_2 . The problem is formulated as hypotheses, with $H_0 : \mu_1 > \mu_2$ and $H_1 : \mu_1 \leq \mu_2$. After n samples, one observes for the parameter vector $\mu = (\mu_1, \mu_2)$ the maximum likelihood estimate $\hat{\mu} = (\hat{\mu}_1, \hat{\mu}_2) = (m_1/n_1, m_2/n_2)$, where m_1 are the successes from distribution 1 after n_1 pulls. Assuming $\hat{\mu}_1 > \hat{\mu}_2$, and assuming that we believe this is the actual state of the world, $\mu = \hat{\mu}$, we assume H_0 and want to test it against the alternative H_1 represented by $\tilde{\mu}$, which is an undefined vector with some $\tilde{\mu}_1 \leq \tilde{\mu}_2$ so that H_1 holds.

Coming from sequential likelihood-ratio tests, Chernoff motivates the use of the Kullback-Leibler divergence to determine from which distribution to sample next. He computes the divergence for the first distribution comparing $\hat{\mu}_1$ and $\tilde{\mu}_1$, as well as the divergence for the second distribution comparing $\hat{\mu}_2$ with $\tilde{\mu}_2$. Chernoff shows to pick the distribution next which maximizes the Kullback-Leibler divergence. The intuition for this comes from how Chernoff computes the alternative $\tilde{\mu}$. It holds that $\tilde{\mu}_1 = \tilde{\mu}_2 = \tilde{\mu}_i$, where $\tilde{\mu}_i$ is a weighted average of $\hat{\mu}_1$ and $\hat{\mu}_2$. As noted by Chernoff (page 759), the weights $(1 - \lambda, \lambda)$ are proportional to how often each distribution has been sampled.

Now, what happens is that long term, we expect the maximum likelihood estimates $\hat{\mu}$ to be close to μ . At the same time, given that $\hat{\mu}$ does not fluctuate much anymore, there will be weights $(1 - \lambda, \lambda)$ such that the two distributions are sampled in fixed proportions (a share of λ for the second distribution), and the Kullback-Leibler divergence for each distribution will be equal: $D_{KL}(\mu_1, \tilde{\mu}_i^*, \nu_1) = D_{KL}(\mu_2, \tilde{\mu}_i^*, \nu_2)$, where $\tilde{\mu}_i^* = (1 - \lambda)\mu_1 + \lambda\mu_2$.

(Is this not just the Chernoff Information as described in Garivier & Kaufmann page 6?)

This observation can be interpreted similar to how under standard composite hypothesis testing, one compares the H_1 against the value under H_0 that is the closest to H_1 . For example, if we observe samples from a normal distribution and estimate the mean with the sample mean which is observed to be positive, $\hat{\mu} > 0$, we might want to test $H_0 : \mu \leq 0$ against $H_1 : \mu > 0$. To do so, we compare $\hat{\mu}$ against the value $\mu_{H_0} \in (\infty, 0]$ closest to H_1 ; thus $\mu_{H_0} = 0$.

In our case, as argued by Chernoff, the mean vector closest to μ under the alternative hypothesis is given by $\tilde{\mu}^*$. Under this value, it is the most difficult to properly distinguish between the different hypotheses. It is the point at which we stand the worst chances to gain new information from sampling any of the two arms (as represented by the fact that both Kullback-Leibler divergences are equal). Consequently (ELABORATE!?), $\tilde{\mu}^* = \arg \min_{\tilde{\mu}} \max_{\nu_i} D_{KL}(\mu, \tilde{\mu}, \nu_i)$.

This result can be described as the solution to a zero-sum two-player minimax game, in which

- the experimenter picks from which distribution to sample from with the goal of maximizing D_{KL}
- then the opponent picks the alternative specification true under the alternative hypothesis to minimize D_{KL}

and the value of the game is given by $\max_{\nu} \min_{\omega} D_{KL}(\mu, \omega, \nu)$. Then, with the opponent picking $\omega = \tilde{\mu}^*$, the experimenter will choose the randomized maxmin strategy of picking ν_1 and ν_2 in proportions $(1 - \lambda)$ and λ .

Very similar games have been described by Russo (2016) and Garivier and Kaufmann (2016). Both analyze settings of pure exploration bandits with fixed confidence. While Garivier and Kaufmann find the game as part of their lower bound on the sample complexity $\mathbb{E}_{\mu}[\tau_{\delta}]$ in their theorem 1, Russo finds his version when deriving the posterior probability on selecting the wrong arm and determines that the solution to the game describes the asymptotically optimal allocation of draws for each arm (page 17).

However, as pointed out by Russo (2016), this game is based on the knowledge of the true parameter μ (θ^* in his notation) and thus the optimal allocation is not readily available when designing an algorithm to solve the exploration problem. Instead, he shows that his algorithms reach the optimal allocation asymptotically even though they do not actively try to solve maxmin game.

(Put in more exact description of the problem solved by Russo so that one can exactly compare the different games and the exponents given by Russo and Garivier & Kaufmann)

In contrast, Garivier and Kaufmann (2016) do devise a sampling rule that is based on estimating the optimal sampling proportions $w^*(\boldsymbol{\mu})$ based on the current estimates $\hat{\boldsymbol{\mu}}(t)$ leading to the plug-in estimates $w^*(\hat{\boldsymbol{\mu}}(t))$. However, the authors acknowledge the fact that using these plug-in estimates in order to estimate the optimal sampling proportions may well lead to problems if the initial estimates $\hat{\boldsymbol{\mu}}(t)$ are off for small t . A bad estimate can lead the bandit algorithm to abandon an arm in the further sampling. Then, an arm that should be explored more – which would lead to a fix of the initial estimate – might not be explored and distort the overall proportions and in the worst case the final decision. In the description of Garivier and Kaufmann, such sampling rule based on plug-in estimates would simply fail from time to time. Thus, they adjust the solution of the optimization problem solving for w^* in order to enforce sufficient exploration of arms, see their lemma 7. More straightforward however is the adjustment in their *D-Tracking* rule in lemma 8, which checks every round whether there exists any arm with pulls $N_a(t) < \sqrt{t} - K/2$ lacking behind the other arms. If this is fulfilled, the arm with the least amount of pulls is pulled. Else, the pull is determined by the plug-in estimates. These rules then also satisfy what was described by Russo: Their empirical sampling proportions converge to the optimal proportions, compare proposition 9 in Garivier and Kaufmann (2016) and proposition 7 in Russo (2016).

0.1.1.1 Game, and Lower Bound on the Sample Complexity

We will now derive the lower bound on the sample complexity in the fixed confidence bandit setting demonstrated in Garivier and Kaufmann (2016).

Consider again the setting as defined in section 0.1. For their lower bound, Garivier and Kaufmann (2016) consider bandit models where the distributions come from one-parameter exponential families. Consequently, the distributions can be parameterized by their means. Therefore, if we let \mathcal{S} to be a set of bandit models with K distributions each, then each model is fully defined by a vector $\boldsymbol{\mu} = (\mu_1, \dots, \mu_K)$. Assume that every $\boldsymbol{\mu} \in \mathcal{S}$ has a unique maximum given by $a^*(\boldsymbol{\mu})$ such that $\mu_{a^*} > \max_{a \neq a^*(\boldsymbol{\mu})} \mu_a$. In order to fix the confidence, we specify the risk δ . The strategies corresponding to this risk are called δ -PAC if for every $\boldsymbol{\mu} \in \mathcal{S}$, $\mathbb{P}_{\boldsymbol{\mu}}(\tau_{\delta} < \infty) = 1$ (the strategy ends surely with after finite samples) and $\mathbb{P}_{\boldsymbol{\mu}}(\hat{a}_{\tau_{\delta}} \neq a^*) \leq \delta$ (the probability of picking not the optimal arm is smaller than the fixed risk). As in Garivier and Kaufmann, we introduce the set of problems for which the optimal arm is different from the one in $\boldsymbol{\mu}$:

$$\text{Alt}(\boldsymbol{\mu}) := \{\boldsymbol{\lambda} \in \mathcal{S} : a^*(\boldsymbol{\lambda}) \neq a^*(\boldsymbol{\mu})\}$$

where $\boldsymbol{\lambda}$ is the vector of means defining the distributions of the alternative problem to $\boldsymbol{\mu}$. Lastly, we have the set of probability distributions on the arms \mathcal{A} : $\Sigma_K = \{w \in \mathbb{R}_+^K : w_1 + \dots + w_K = 1\}$. We are now ready to state the theorem.

Theorem 1 (Garivier and Kaufmann, 2016) *Let $\delta \in (0, 1)$. For any δ -PAC strategy and any bandit model $\boldsymbol{\mu} \in \mathcal{S}$,*

$$\mathbb{E}_{\boldsymbol{\mu}}[\tau_{\delta}] \geq T^*(\boldsymbol{\mu}) \text{kl}(\delta, 1 - \delta),$$

where

$$T^*(\boldsymbol{\mu})^{-1} := \sup_{w \in \Sigma_K} \inf_{\boldsymbol{\lambda} \in \text{Alt}(\boldsymbol{\mu})} \left(\sum_{a=1}^K w_a d(\mu_a, \lambda_a) \right). \quad (1)$$

In equation (1), we again see the game occur that we describe above. In this case, the game is about the minimal number of rounds $T = T(w, \lambda)$ needed for any bandit strategy to ensure that the strategy fulfills the δ -PAC constraint. Consider that we have the inverse of T in equation (1). Then we see that

- the statistician tries to minimize the number of rounds by choosing the optimal proportions of draws w^* for each arm
- the opponent picks the alternative model $\boldsymbol{\lambda}$ so as to maximize the number of necessary rounds

When choosing $\boldsymbol{\lambda}$, the opponent will choose an alternative as close to the original model as possible so as to minimize the Kullback-Leibler divergence $d(\mu_a, \lambda_a)$. As already noted in Chernoff (1959), this will minimize the information available to the statistician and make it difficult to distinguish between the cases with confidence.




Figure 1: Statistician chooses the curve by picking the weights w , then the opponent picks λ . In the case of a two-armed bandit, we have $w_1 + w_2 = 1$. The plot suggests that uniform allocation is optimal for Bernoulli bandits. The thick line shows the result for $w_1 = w_2 = 0.5$.

The plot above specifies again the game and the reasoning of the opponent. Given a model $\boldsymbol{\mu} = (\mu_1, \dots, \mu_K)$ with $\mu_1 > \mu_2 \geq \dots \geq \mu_K$, the experimenter chooses the proportions of arm draws $w = (w_a)_a$. Then the opponent chooses an arm $a \in \{2, \dots, K\}$ and chooses the

alternative distribution for arm a , $\lambda_a = \arg \min_{\lambda} w_1 d(\mu_1, \lambda) + w_a d(\mu_a, \lambda)$. The payoff of the game is then the minimal number of rounds T so that the restriction $T w_1 d(\mu_1, \lambda_a - \epsilon) + T w_a d(\mu_a, \lambda + \epsilon) \geq \text{kl}(\delta, 1 - \delta)$ holds when $a^*(\boldsymbol{\mu}) \neq a^*(\boldsymbol{\lambda})$ due to ϵ . Consequently, the value is given by

$$T(w, a, \delta) = \frac{\text{kl}(\delta, 1 - \delta)}{w_1 d(\mu_1, \lambda_a - \epsilon) + w_a d(\mu_a, \lambda + \epsilon)} \quad (2)$$

which is comparable to the result in Theorem (Garivier and Kaufmann, 2016) [1](#).

0.1.2 Fixed Budget